
Transactional Memory and concurrency

Simon Peyton Jones
Microsoft Research, Cambridge

The Context

Multicore



**Parallel
programming
essential**



Task parallelism

- Explicit threads
- Synchronise via locks, messages

This talk

Data parallelism

Operate simultaneously on bulk data

1.30pm today...

Task parallelism: state of play

- The state of the art in concurrent programming is 30 years old: locks and condition variables. (In Java/C#: synchronised methods, lock statements, Monitor.wait.)
- Locks and condition variables are fundamentally flawed: it's like building a sky-scraper out of bananas.
- **This presentation describes significant recent progress: bricks and mortar instead of bananas**


What's wrong with locks?

A 30-second review:



- **Races:** due to forgotten locks
- **Deadlock:** locks acquired in “wrong” order.
- **Lost wakeups:** forgotten notify to condition variable
- **Diabolical error recovery:** need to restore invariants and release locks in exception handlers
- These are serious problems. But even worse...

Locks are absurdly hard to get right

Scalable double-ended queue: one lock per cell



No interference if
ends "far enough"
apart



But watch out when the queue
is 0, 1, or 2 elements long!

Locks are absurdly hard to get right

Coding style	Difficulty of concurrent queue
Sequential code	Undergraduate

Locks are absurdly hard to get right

Coding style	Difficulty of concurrent queue
Sequential code	Undergraduate
Locks and condition variables	Publishable result at international conference

Atomic memory transactions

Coding style	Difficulty of concurrent queue
Sequential code	Undergraduate
Locks and condition variables	Publishable result at international conference
Atomic blocks	Undergraduate

Atomic memory transactions

Like
database
transactions

```
atomic { ... sequential get code ... }
```

- To a first approximation, just write the sequential code, and wrap **atomic** around it
- All-or-nothing semantics: **Atomic** commit
- Atomic block executes in **Isolation**
- Cannot deadlock (there are no locks!)
- Atomicity makes error recovery easy (e.g. exception thrown inside the **get** code)

AcId

How does it work?

Optimistic concurrency

```
atomic { ... <code> ... }
```

One possibility:

- Execute <code> without taking any locks
- Each read and write in <code> is logged to a thread-local transaction log
- Writes go to the log only, not to memory
- At the end, the transaction tries to **commit** to memory
- Commit may fail; then transaction is re-run

Blocking

```
atomic { if n_items == 0 then retry  
        else ...remove from queue... }
```

- **retry** says “abandon the current transaction and re-execute it from scratch”
- The implementation waits until `n_items` changes
- No condition variables, no lost wake-ups!

Blocking composes

```
atomic { x = queue1.getItem()  
        ; queue2.putItem( x ) }
```

- If either **getItem** or **putItem** retries, the whole transaction retries
- So the transaction waits until **queue1** is not empty AND **queue2** is not full
- No need to re-code **getItem** or **putItem**
- (Lock-based code does not compose)

Choice

```
atomic { x = queue1.getItem()
        ; choose
          queue2.putItem(x)
        orElse
          queue3.putItem(x) }
```

- **orElse** tries two alternative paths
- If the first retries, it runs the second
- If both retry, the whole **orElse** retries.

Choice composes too

```
atomic { x = queue1.getItem()  
        ; choose  
            queue2.putItem(x)  
        orElse  
            queue3.putItem(x) }
```


- So the transaction waits until
 - queue1 is non-empty, AND
 - EITHER queue2 is not full OR queue3 is not fullwithout touching **getItem** or **putItem**

Performance


There's a run-time cost for TM, but

- Compiler technology and hardware support can reduce it a lot
- A “faster” program that doesn't work right is useless
- TM allows much finer-grain locking without losing correctness, so performance may be **better** than when using locks

Results: concurrency control overhead



Results: scalability



Remember this

- Atomic blocks (atomic, retry, orElse) are a real step forward
- **It's like using a high-level language instead of assembly code:** whole classes of low-level errors are eliminated.
- Not a silver bullet:
 - you can still write buggy programs;
 - concurrent programs are still harder to write than sequential ones;
 - aimed at shared memory
- Very hot research area – expect developments
- Available in STM Haskell today: <http://haskell.org/ghc>

Backup slides

Starvation

- A worry: could the system “thrash” by continually colliding and re-executing?
- No: one transaction can be forced to re-execute only if another succeeds in committing. That gives a strong progress guarantee.
- But a particular thread could perhaps starve.
- No automatic solution can possibly be adequate


No I/O inside transactions

```
atomic { if ( $x > y$ ) then launchMissiles }
```

- The transaction might see ($x > y$) because it pauses between reading x and reading y
- So we must not call **launchMissiles** until the transaction commits
- Simple story: no I/O inside transactions

Input/output

- Transactional output is easy:



- Input is a bit harder, because of the need to make sure the transactional input buffer is filled enough